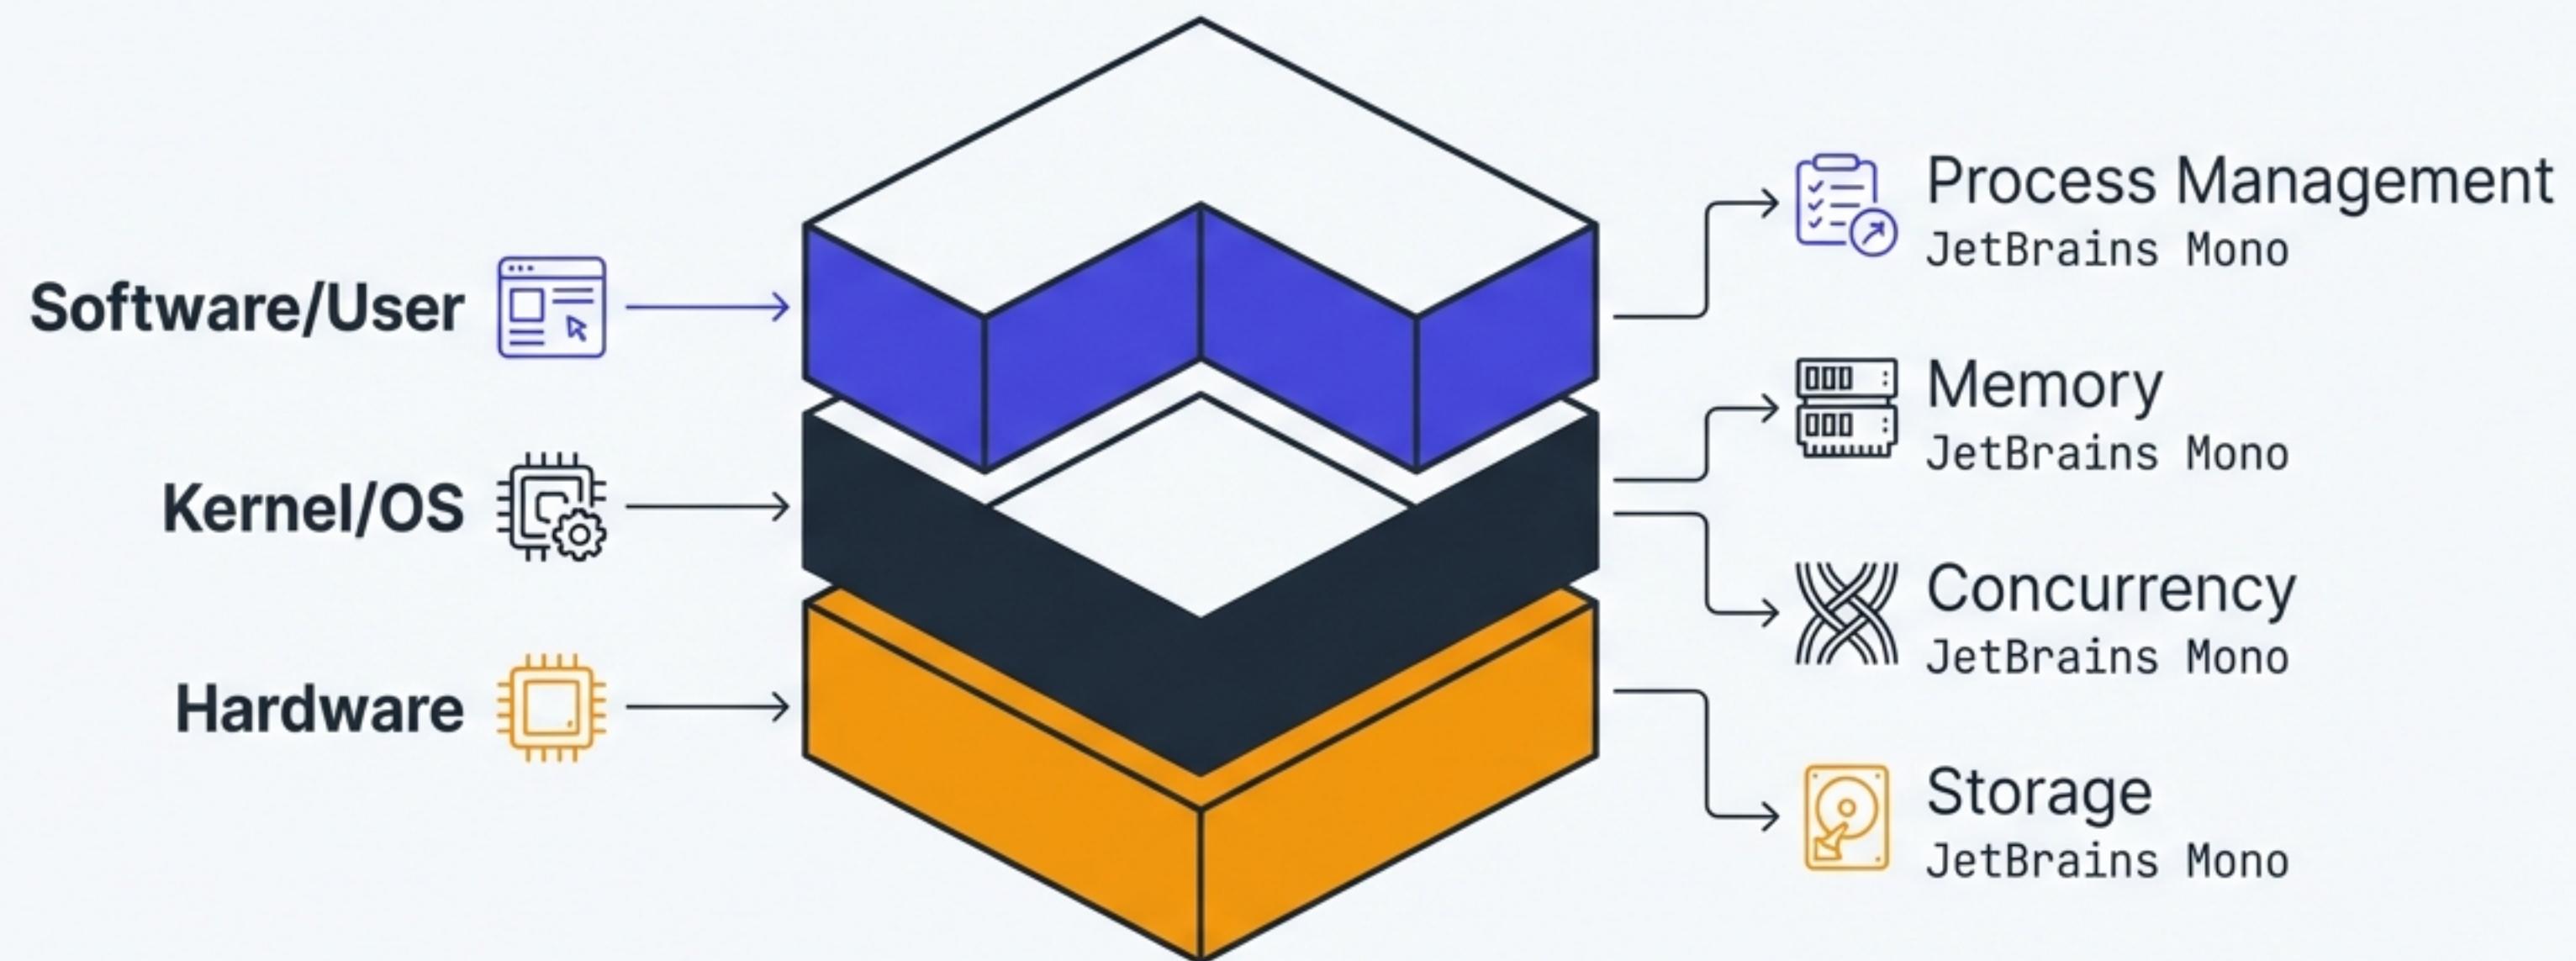


# Operating Systems Masterclass

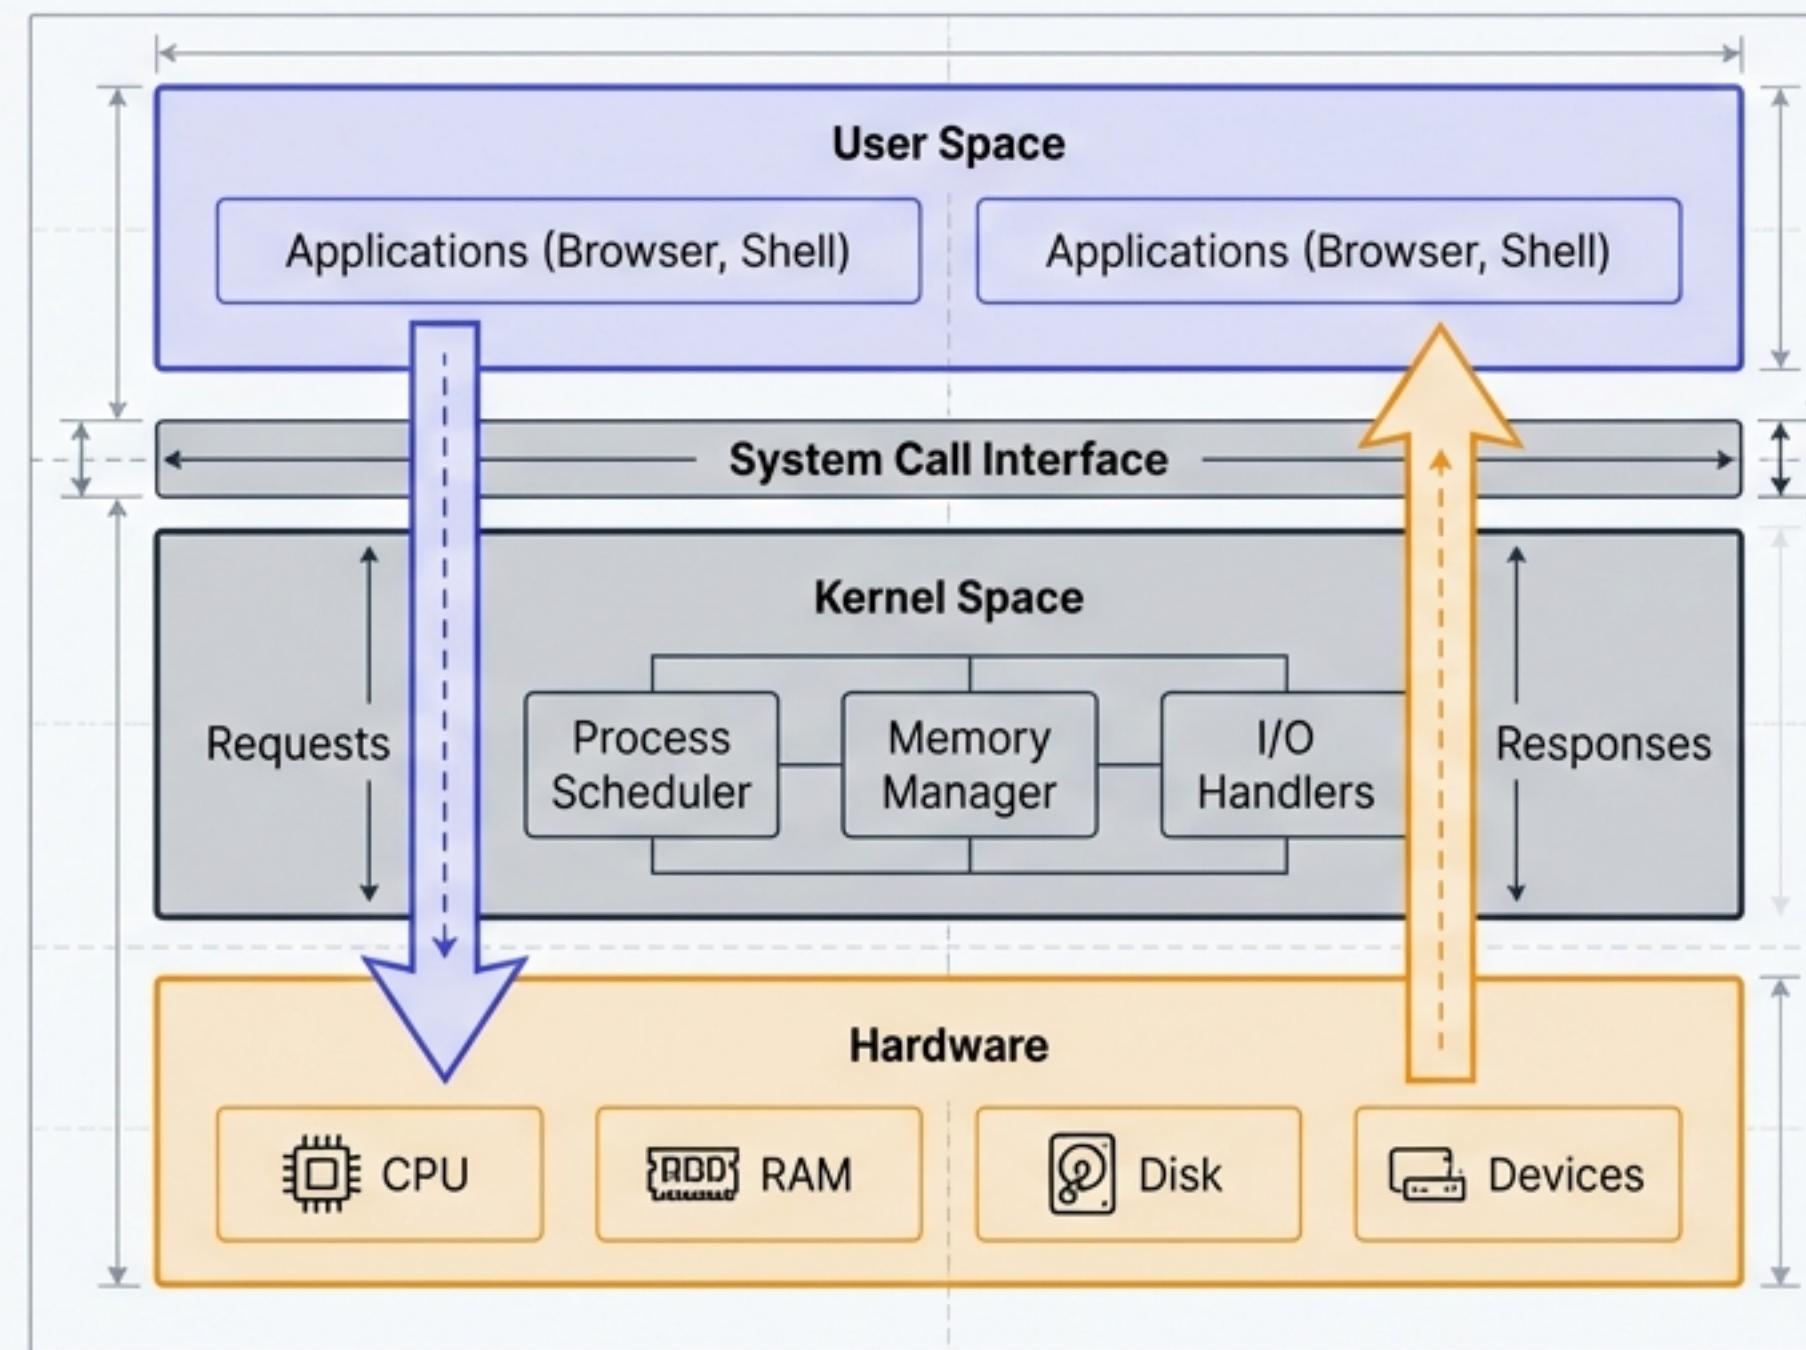
## Architecture, Internals & System Design



COMPREHENSIVE REVIEW: CORE • INTERMEDIATE • ADVANCED

# The Role of the Conductor: OS & Kernel

Mediating between user intent and hardware reality



## OS vs. Kernel

The **OS** includes the full suite of system software (GUI, Shell, Tools).

The **Kernel** is the strictly core component bridging software to hardware.

## Kernel Types

**Monolithic:** User and Kernel services in same address space (Faster, Larger).

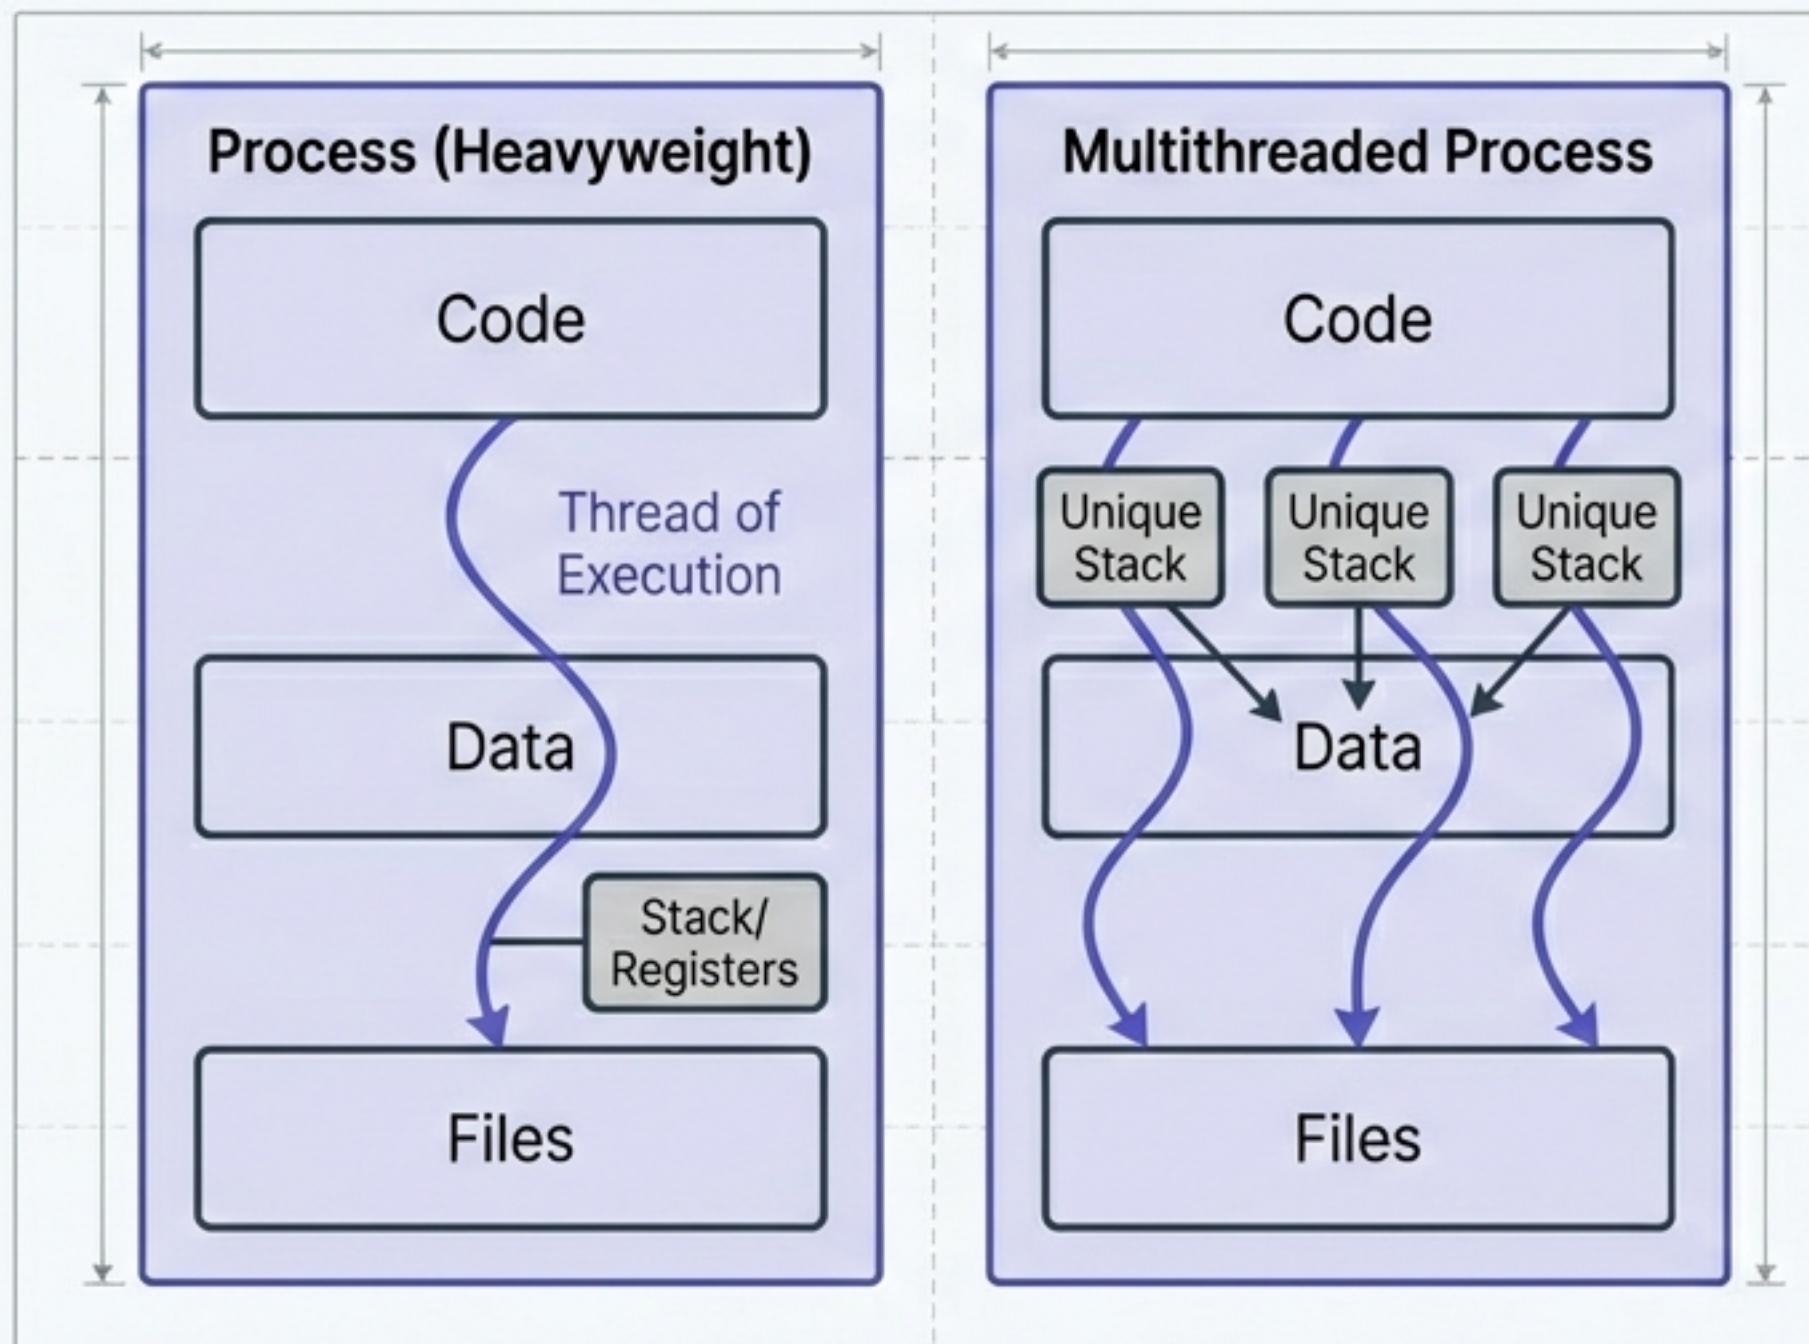
**Microkernel:** Minimal services in kernel mode (Modular, Slower).

## The Bootstrap

The initial program loaded by the POST (Power-On Self-Test) to initialize the OS components.

# Units of Execution: Process vs. Thread

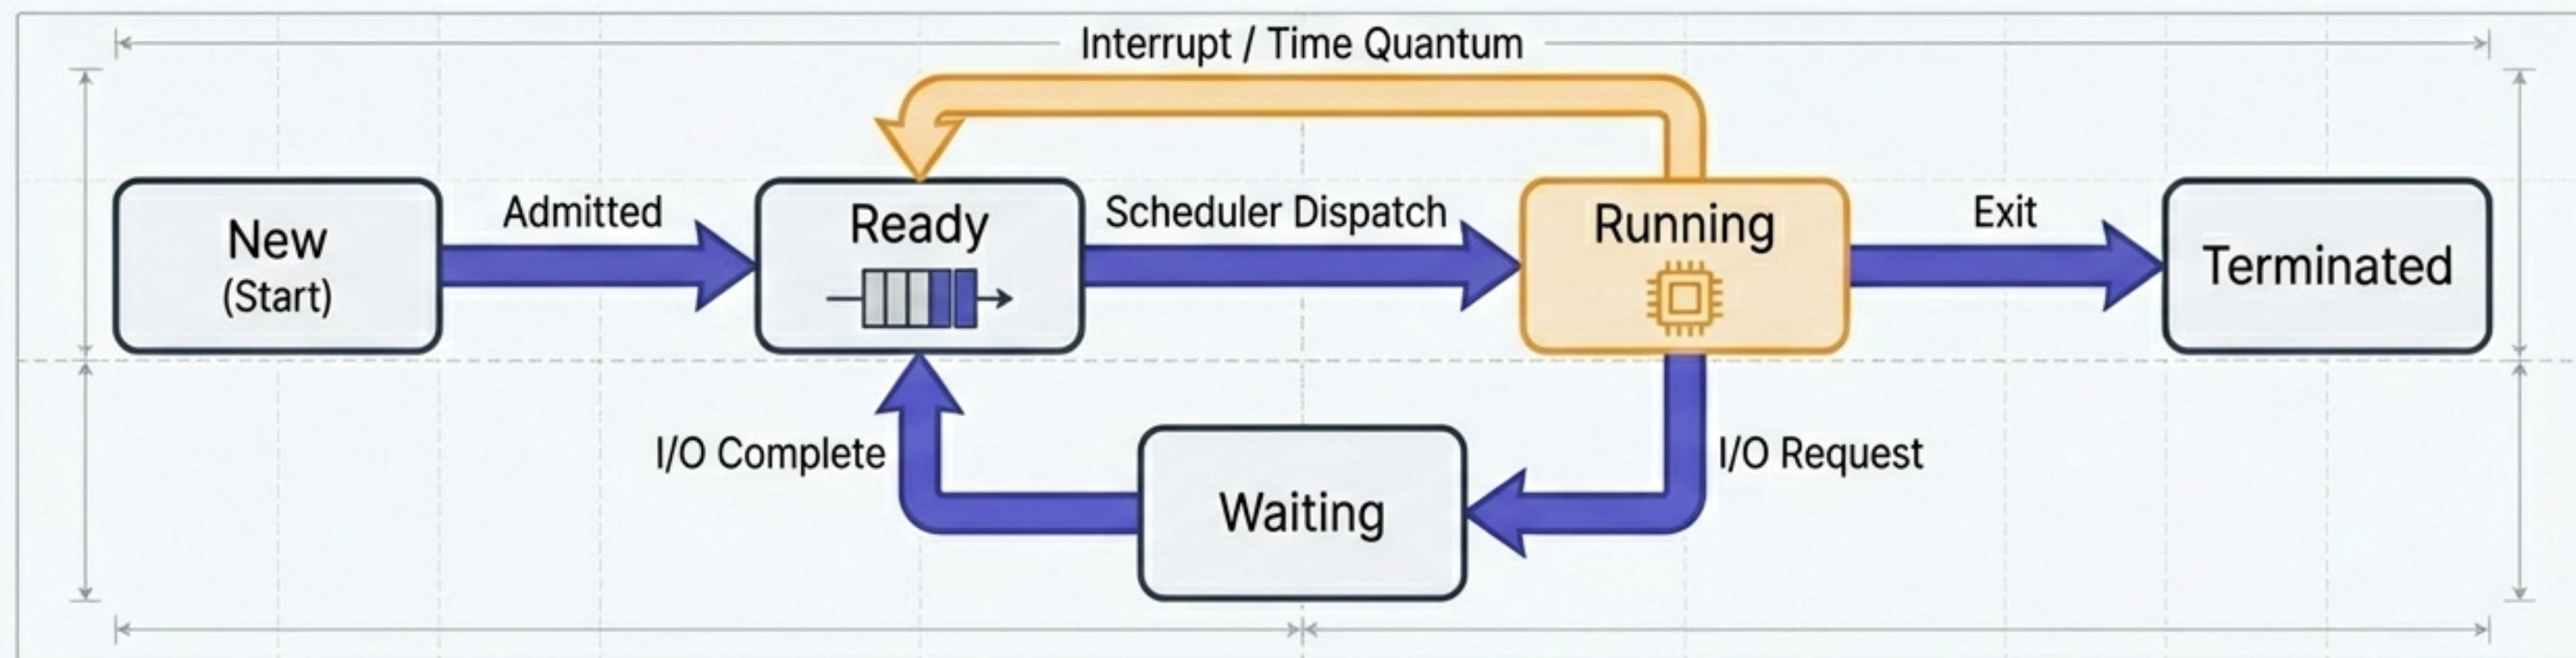
Heavyweight isolation versus lightweight concurrency.



Feature	Process	Thread
Memory	Isolated Address Space	Shared Memory (Heap/Data)
Creation Cost	High (OS allocates resources)	Low (Lightweight)
Communication	Slow (IPC mechanisms)	Fast (Direct shared access)
Context Switch	Heavy (Save full state)	Light (Save registers/stack)
Failure Scope	Isolated (Others run)	Fatal (Can crash process)

# The Process Lifecycle & State Management

Understanding process transitions, data structures, and system anomalies.



## PCB (Process Control Block)

The data structure tracking execution status. Contains Program Counter, Registers, and Priority. Stored in the Process Table.

## Context Switching

The overhead of saving the state of the old process (to PCB) and loading the new one. No useful work is done during this switch.

## Lifecycle Anomalies

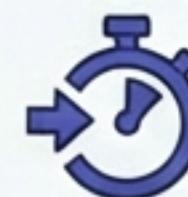
**Zombie:** Finished execution but entry remains in process table (Parent hasn't read exit code).

**Orphan:** Parent terminated while child is still running.

# Scheduling the CPU: Algorithms

Deciding which process runs next to maximize throughput. in Inter.

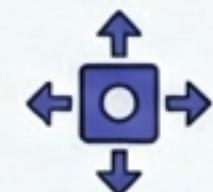
## Scheduling Strategies



**"Preemptive"**: OS allocates CPU for limited time (**Quantum**). Higher priority can interrupt. (Flexible, High Overhead).



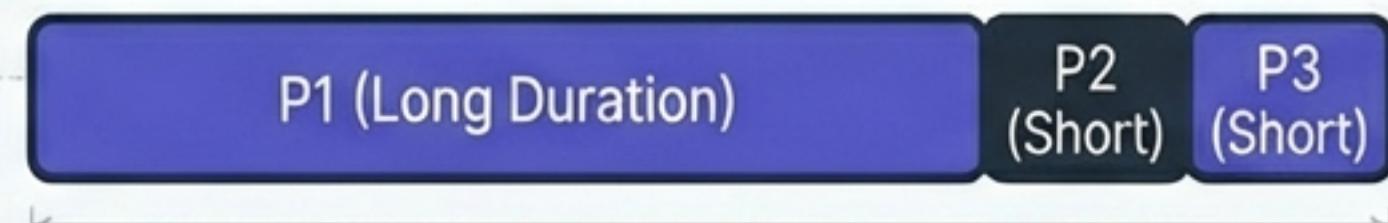
**"Non-Preemptive"**: Process holds CPU until termination or I/O wait. (Simple, Rigid).



**"Dispatcher"**: The module that gives control of the CPU to the selected process.

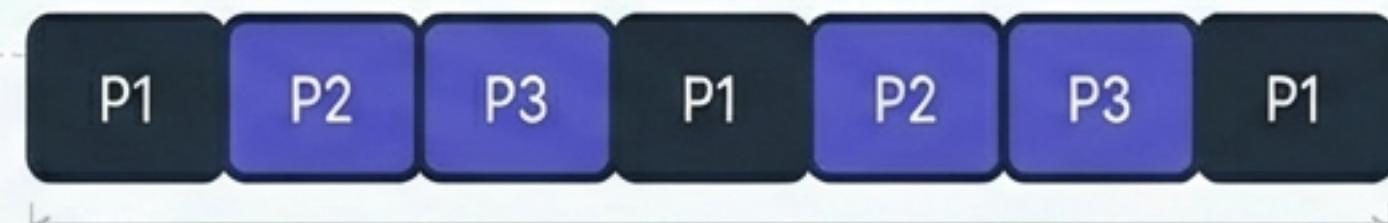
## Algorithm Comparison

### FCFS (First-Come, First-Served)



Convoy Effect: Short jobs wait for long job.

### Round Robin (RR)

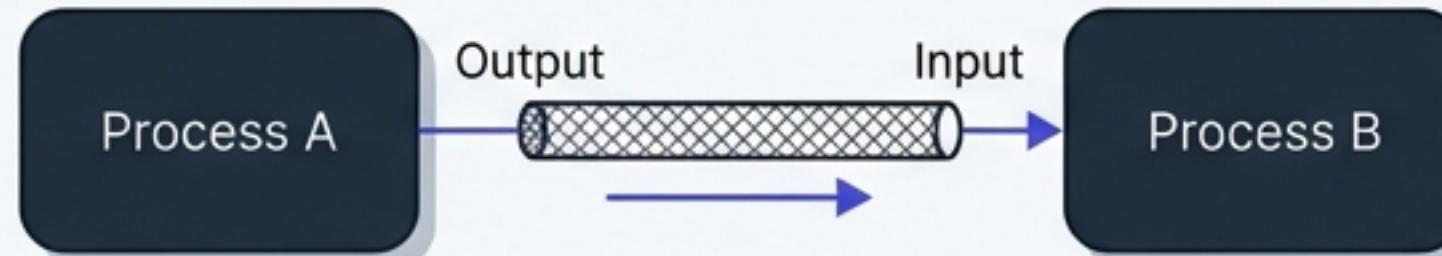


Time Quantum (Time Slicing). Fair, but higher context switching overhead.

# Inter-Process Communication (IPC)

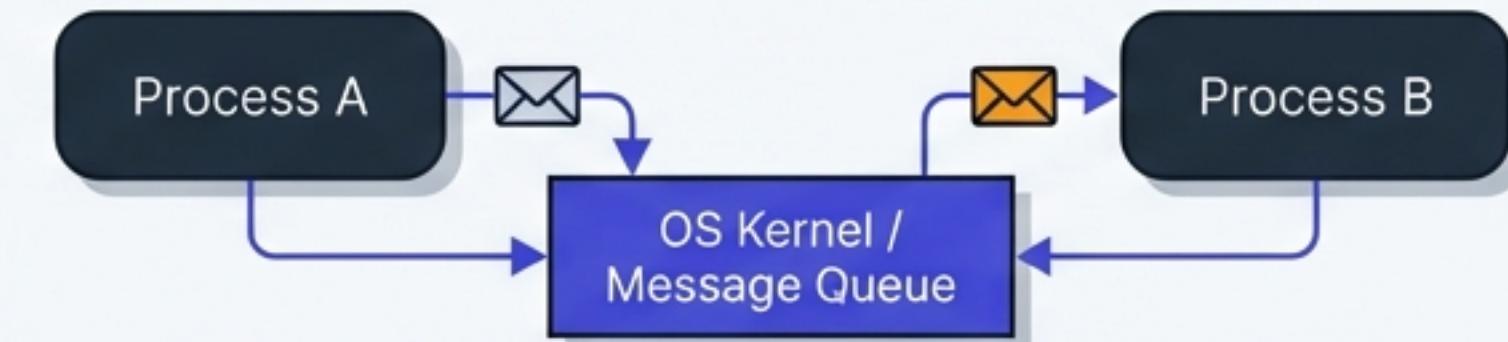
Mechanisms for isolated processes to exchange data.

## Pipes



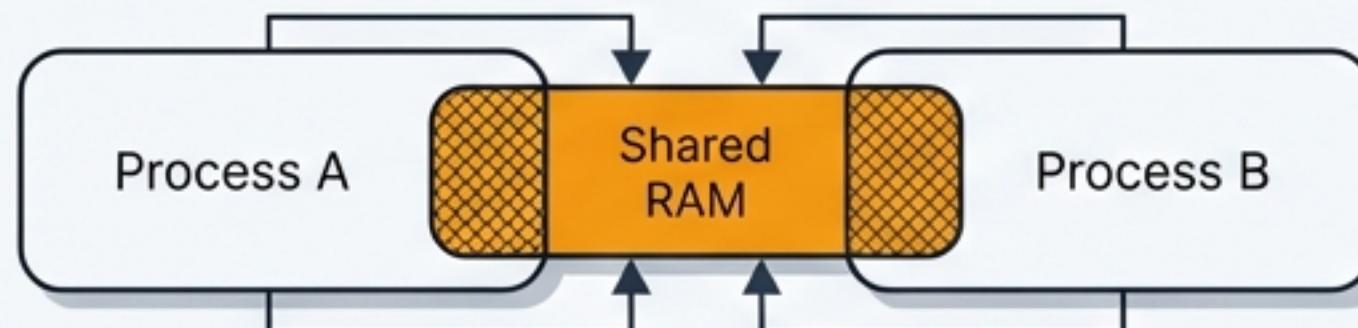
Unidirectional data flow.  
Output of one becomes input of another.

## Message Queue



Kernel-managed list of messages  
Asynchronous communication.

## Shared Memory



Fastest method. Direct memory access.  
Requires synchronization to prevent conflicts.

## Sockets



Network-based communication.  
Works across different machines/OS.

# The Synchronization Challenge

Race Conditions and the Critical Section Problem.

## Race Condition Visualization



When two threads access a shared resource (variable/file) simultaneously, the outcome depends on execution order.



### Mutex (Lock)

A key. Only the holder can enter.  
If locked, others wait.  
Prevents simultaneous access.



### Semaphore

A signaling integer. Non-negative. `Wait()` decrements,  
`Signal()` increments.  
Can allow N simultaneous users  
(Counting Semaphore).



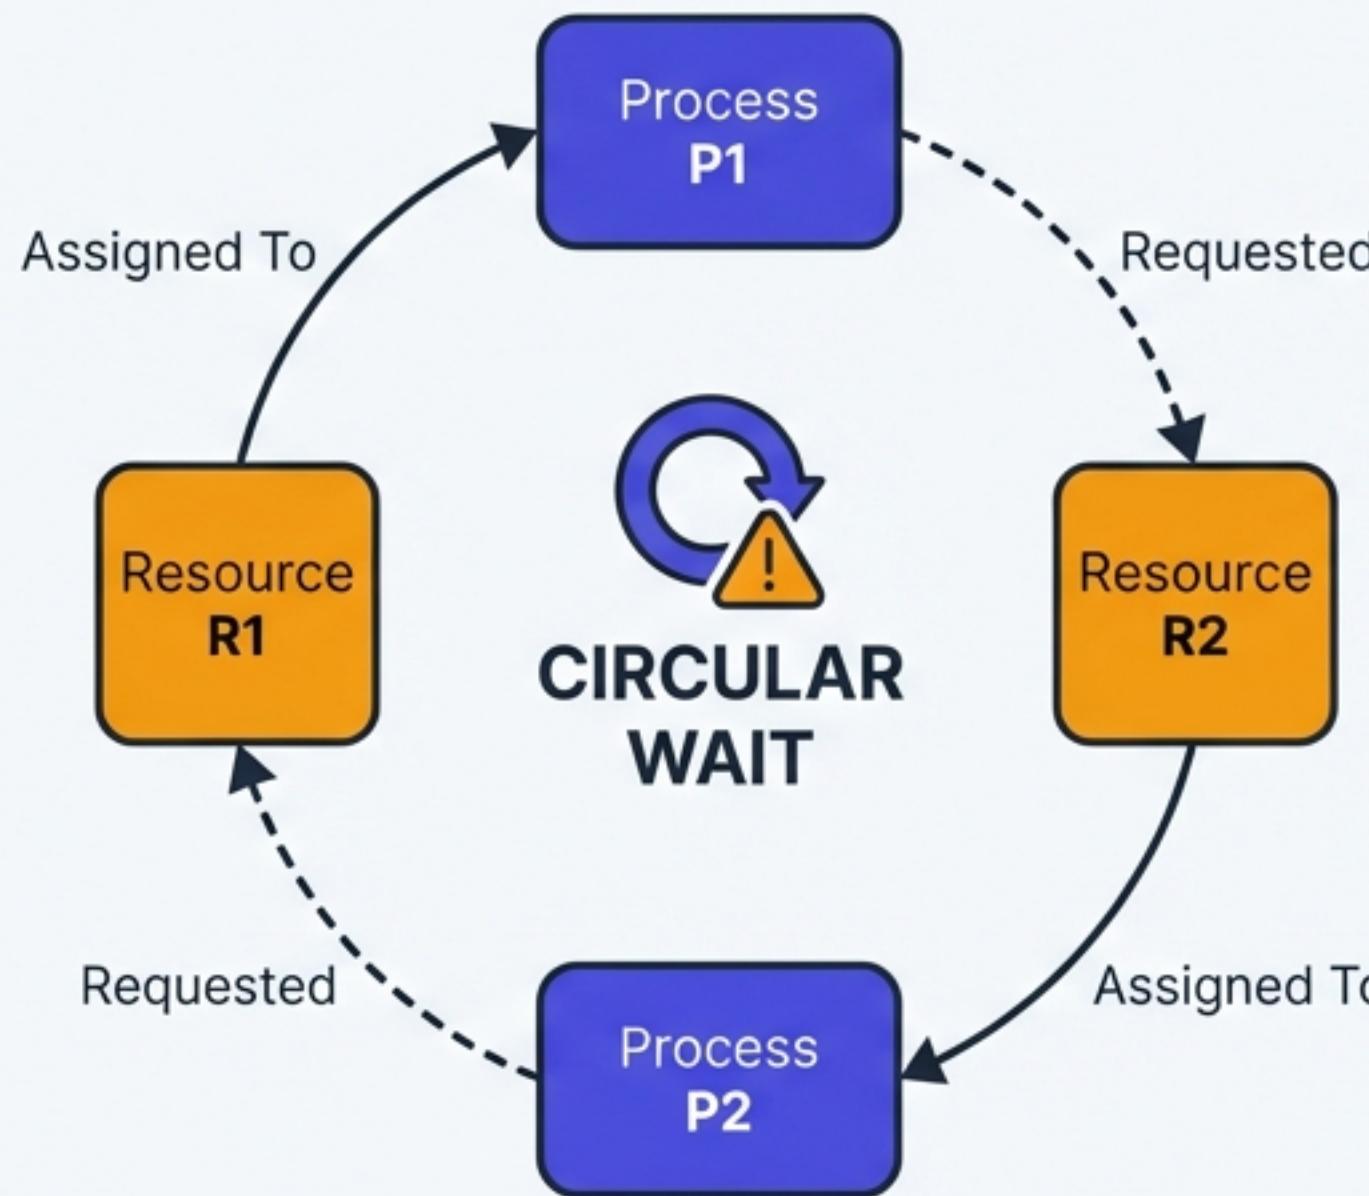
### Peterson's Solution

Software-based algorithmic approach using two variables (flag and turn) to enforce mutual exclusion.

# The Nightmare Scenario: Deadlock

A state where processes wait indefinitely for each other.

## Resource Allocation Graph



## The 4 Necessary Conditions (Coffman Conditions)

If one is broken, deadlock is impossible.



### Mutual Exclusion

Resources cannot be shared (e.g., a printer).



### Hold and Wait

Holding a resource while waiting for another.



### No Preemption

Resources cannot be forcibly taken away.

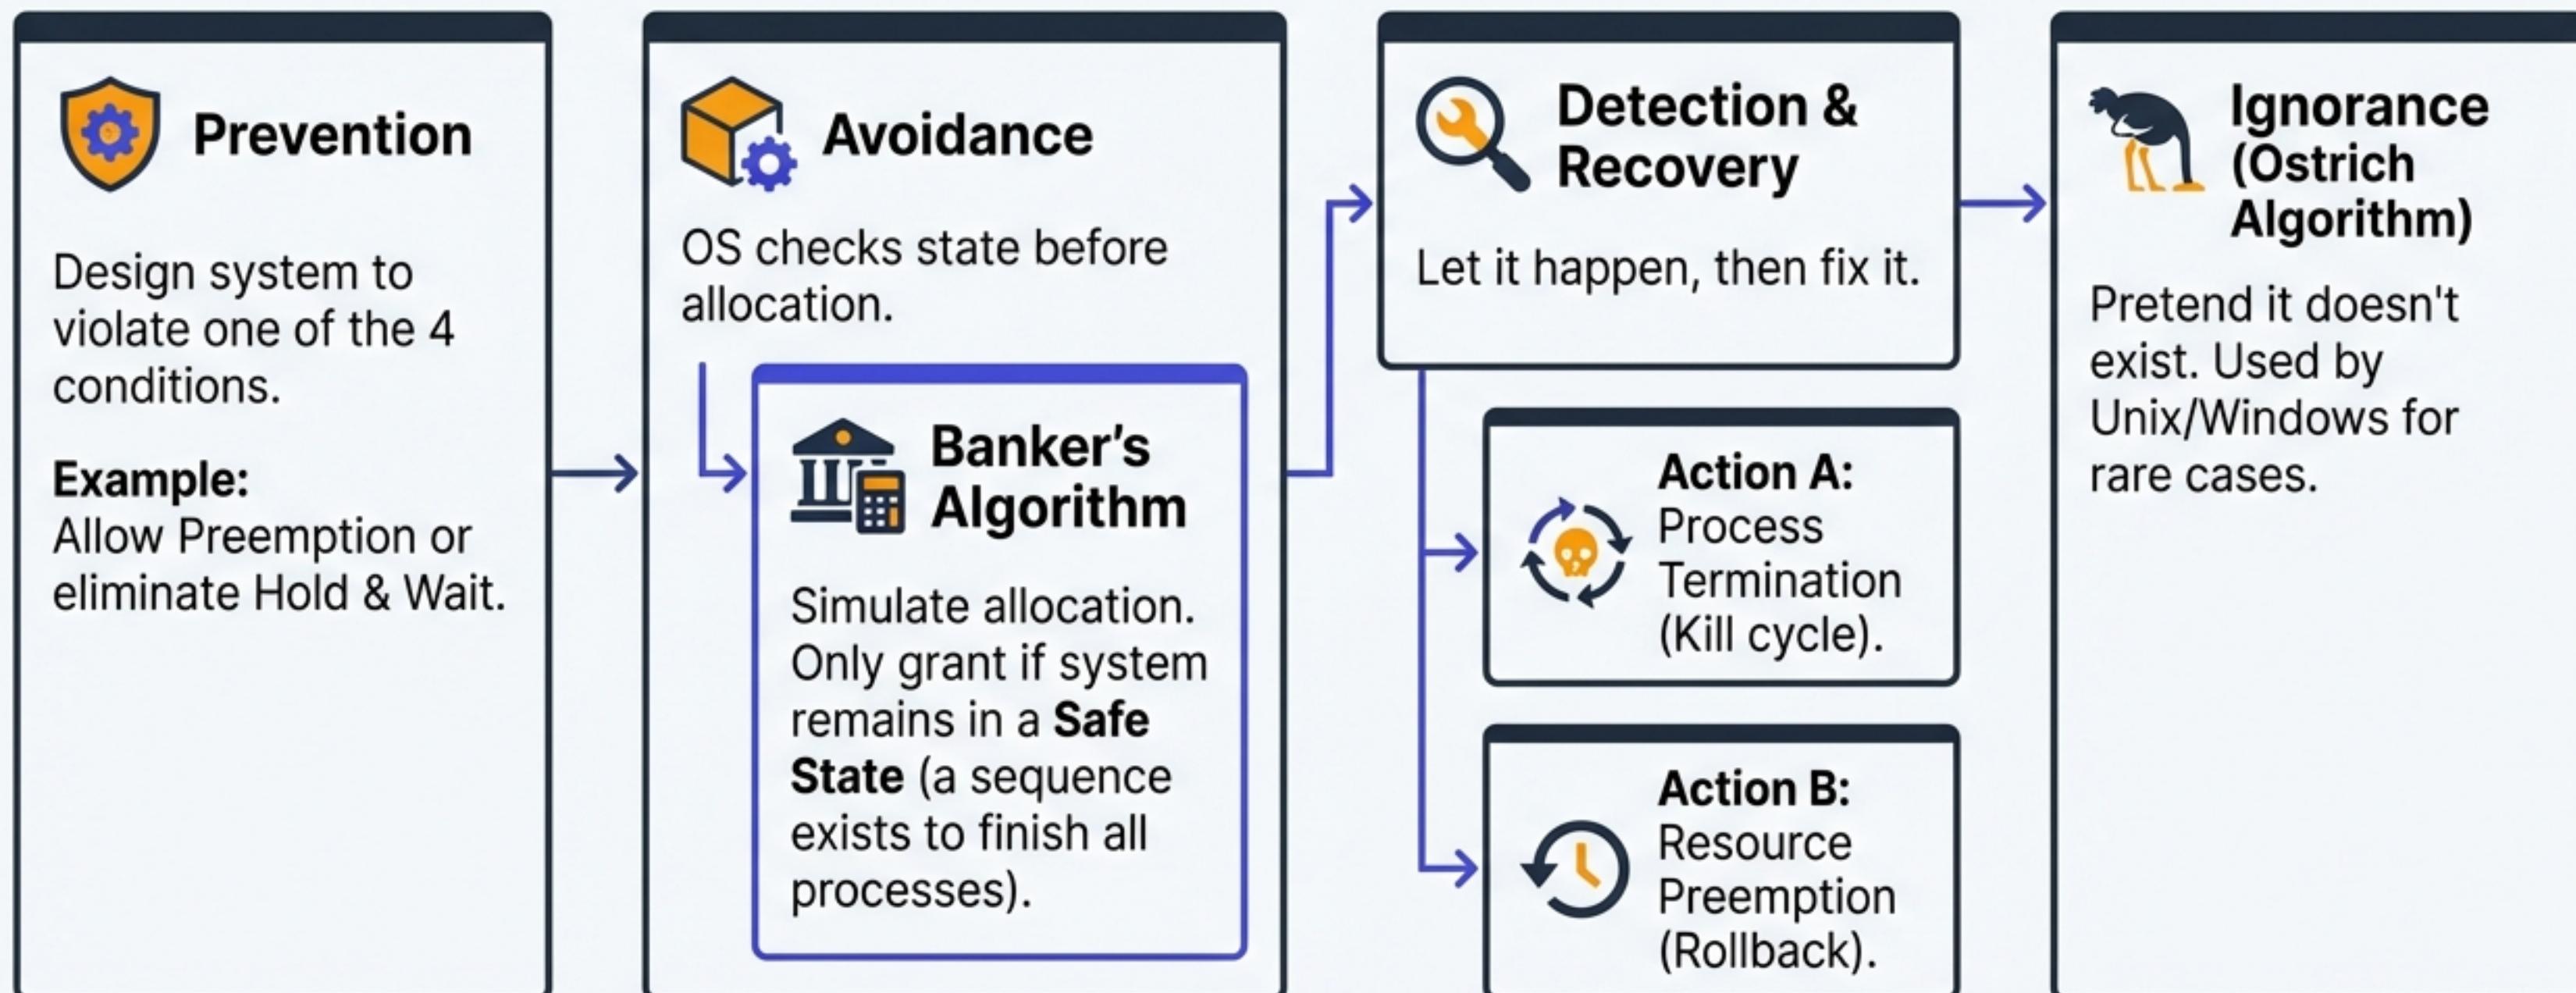


### Circular Wait

A closed chain of processes waiting for each other.

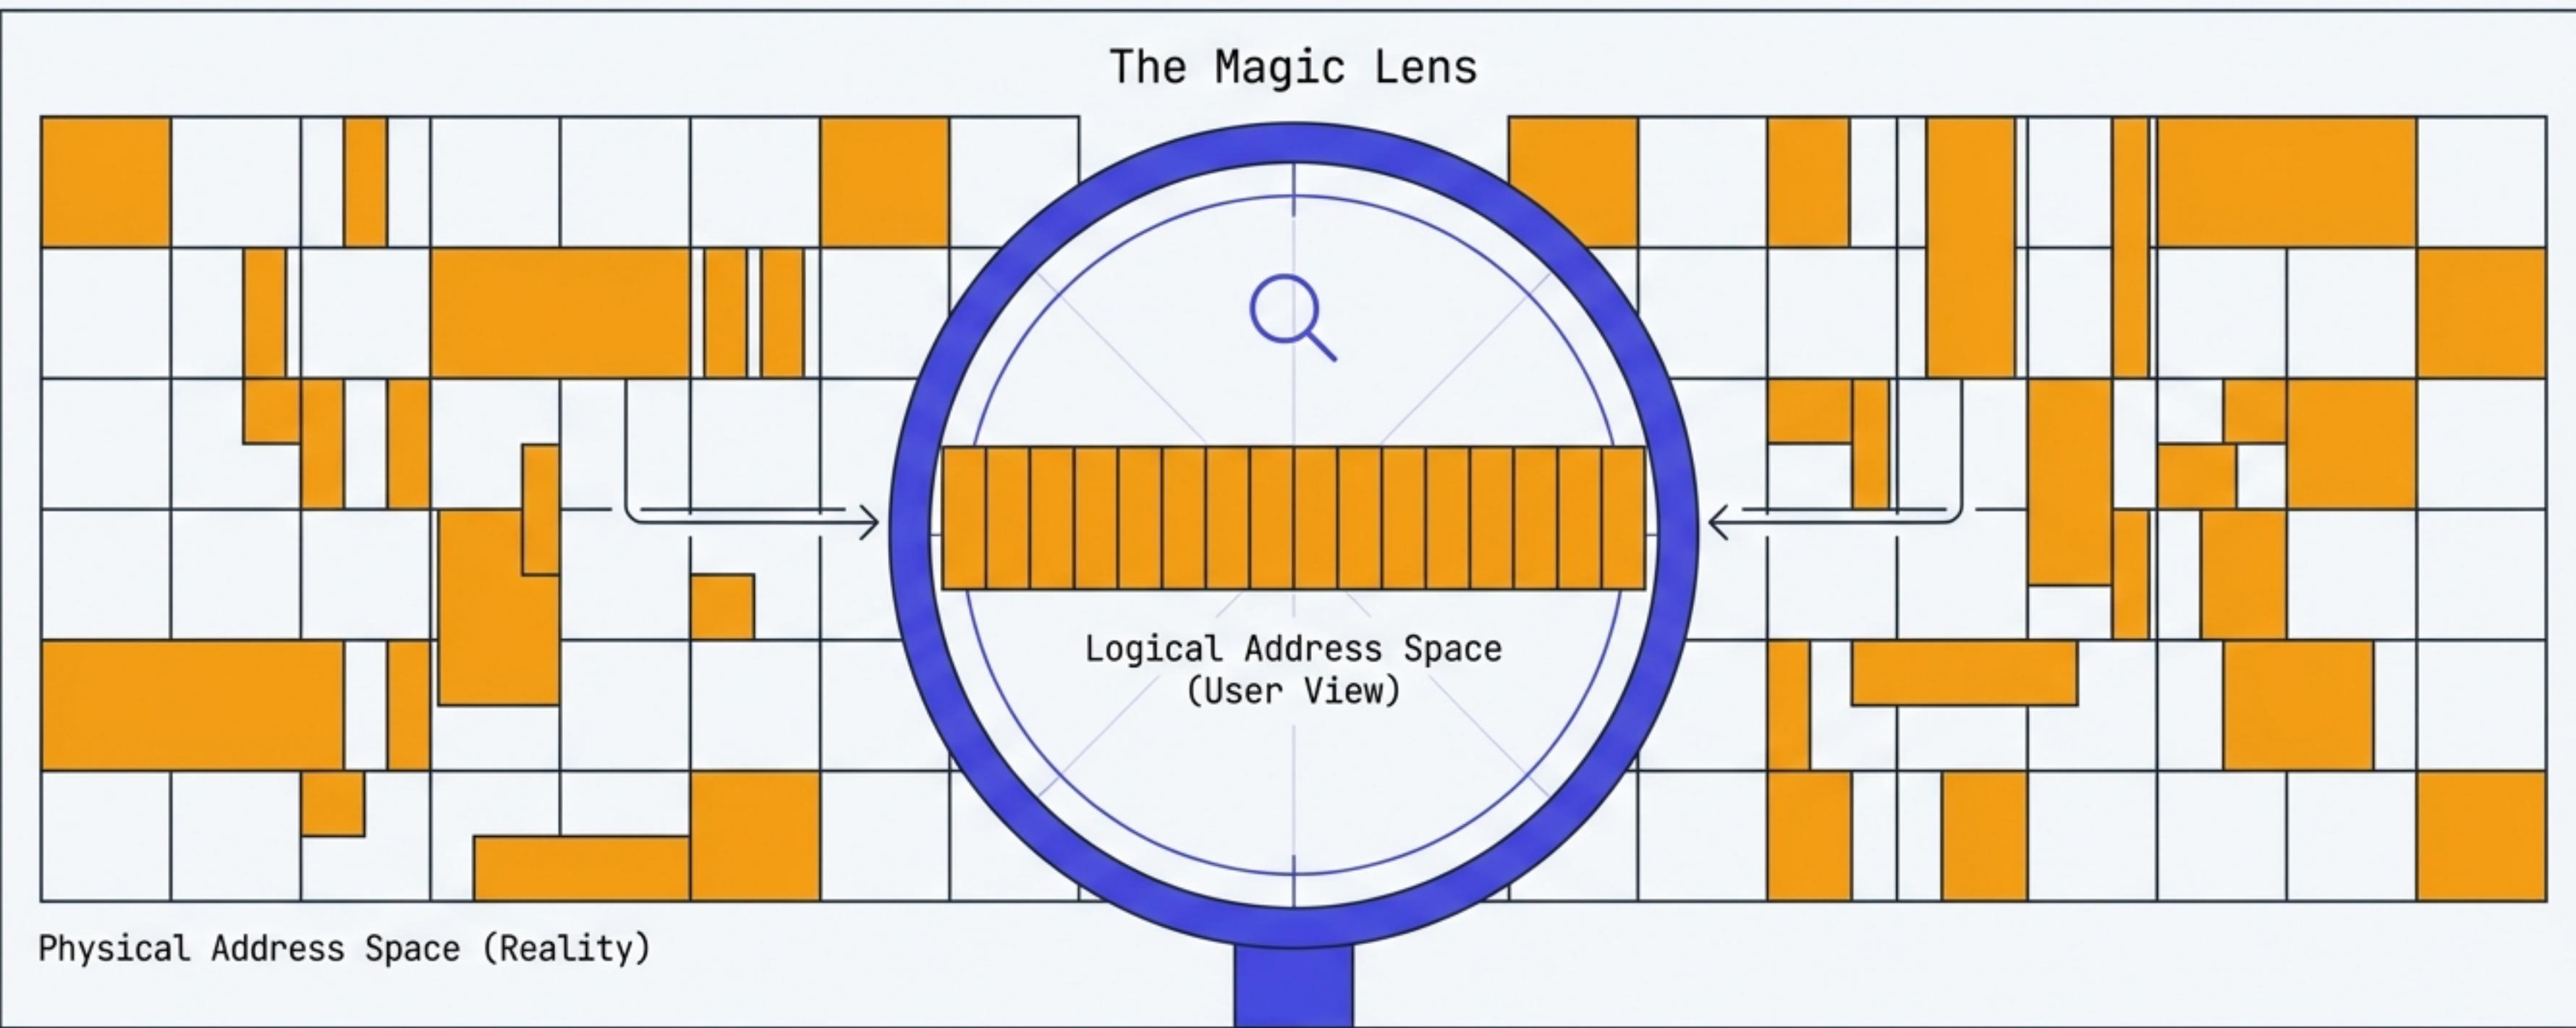
# Handling Deadlocks

Strategies for Prevention, Avoidance, and Recovery.



# Memory Abstraction: Virtual vs. Physical

The illusion of infinite, contiguous memory.

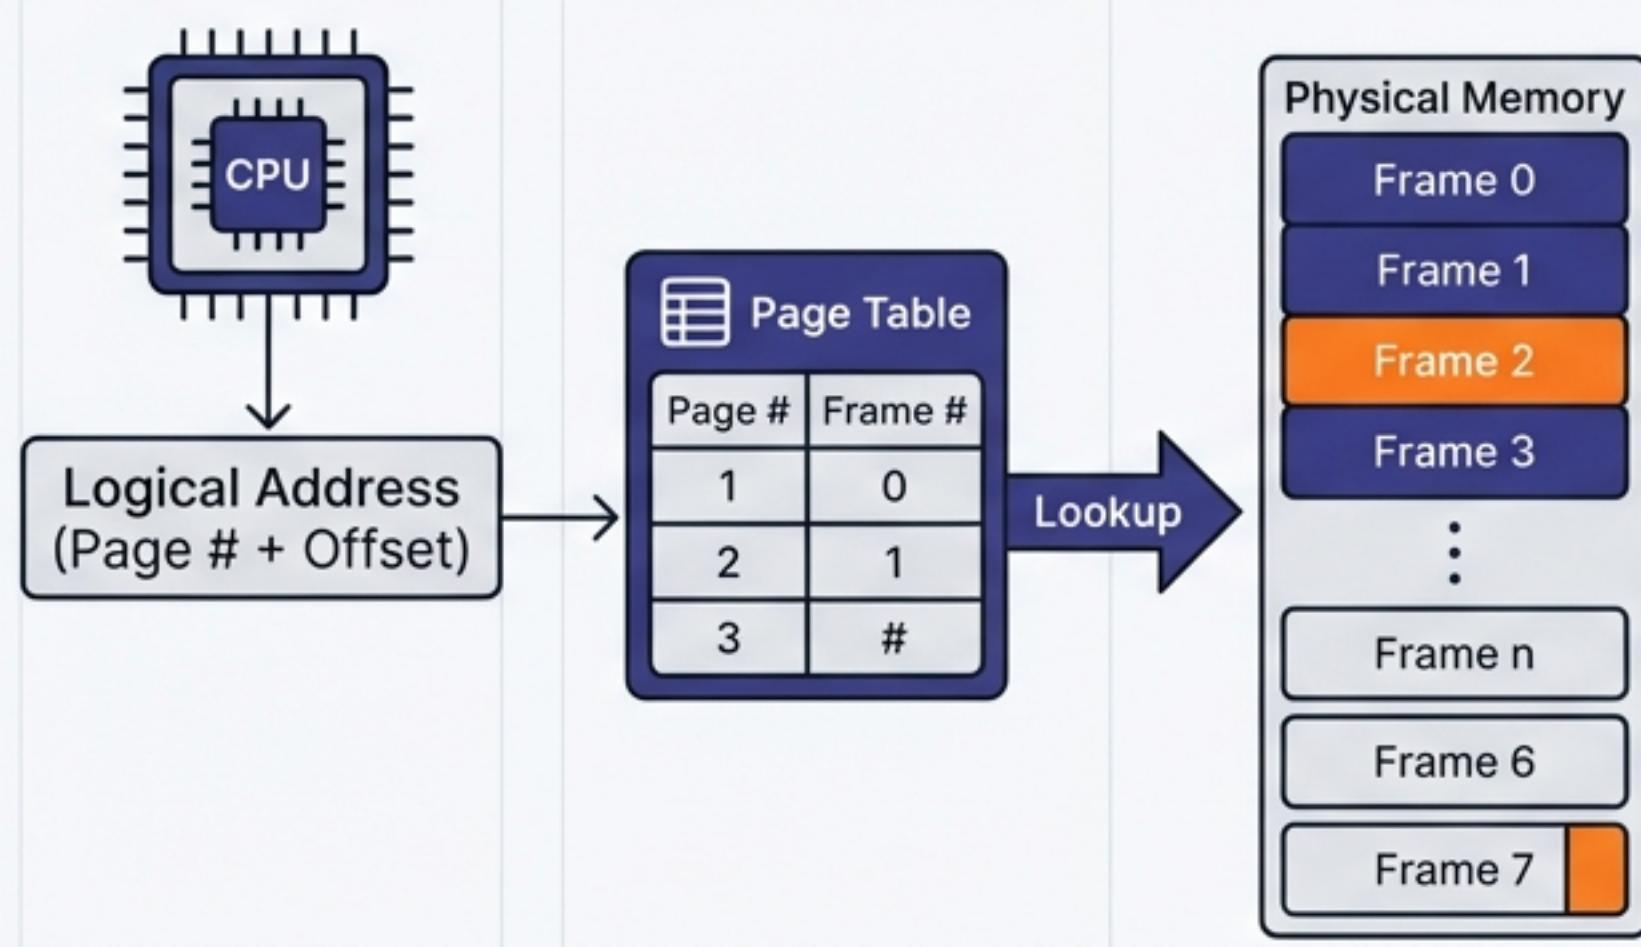


**Memory Management Unit (MMU):** Hardware device that translates Logical Addresses (CPU generated) to Physical Addresses (RAM location) at runtime.

**Benefit:** Allows execution of programs larger than physical memory and provides process isolation.

# Organizing Memory: Paging & Segmentation

## Paging (Fixed Size)

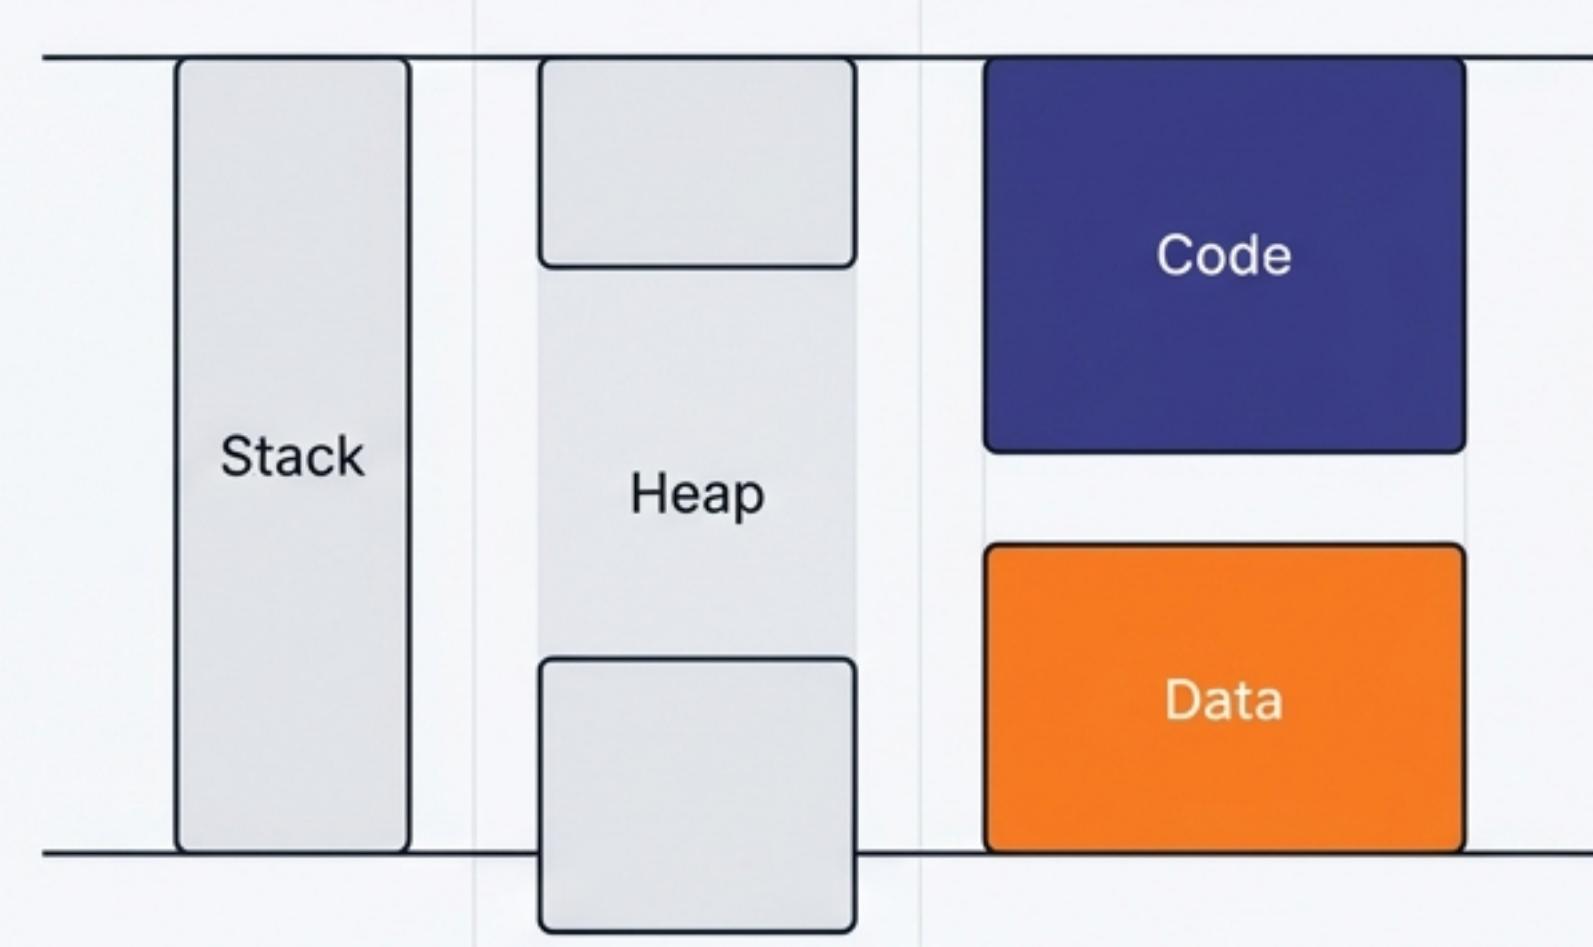


Invisible to user. Solves external fragmentation but causes internal fragmentation (wasted space inside last page).



\*\*Demand Paging\*\*: Loading pages only when a 'Page Fault' occurs.

## Segmentation (Variable Size)

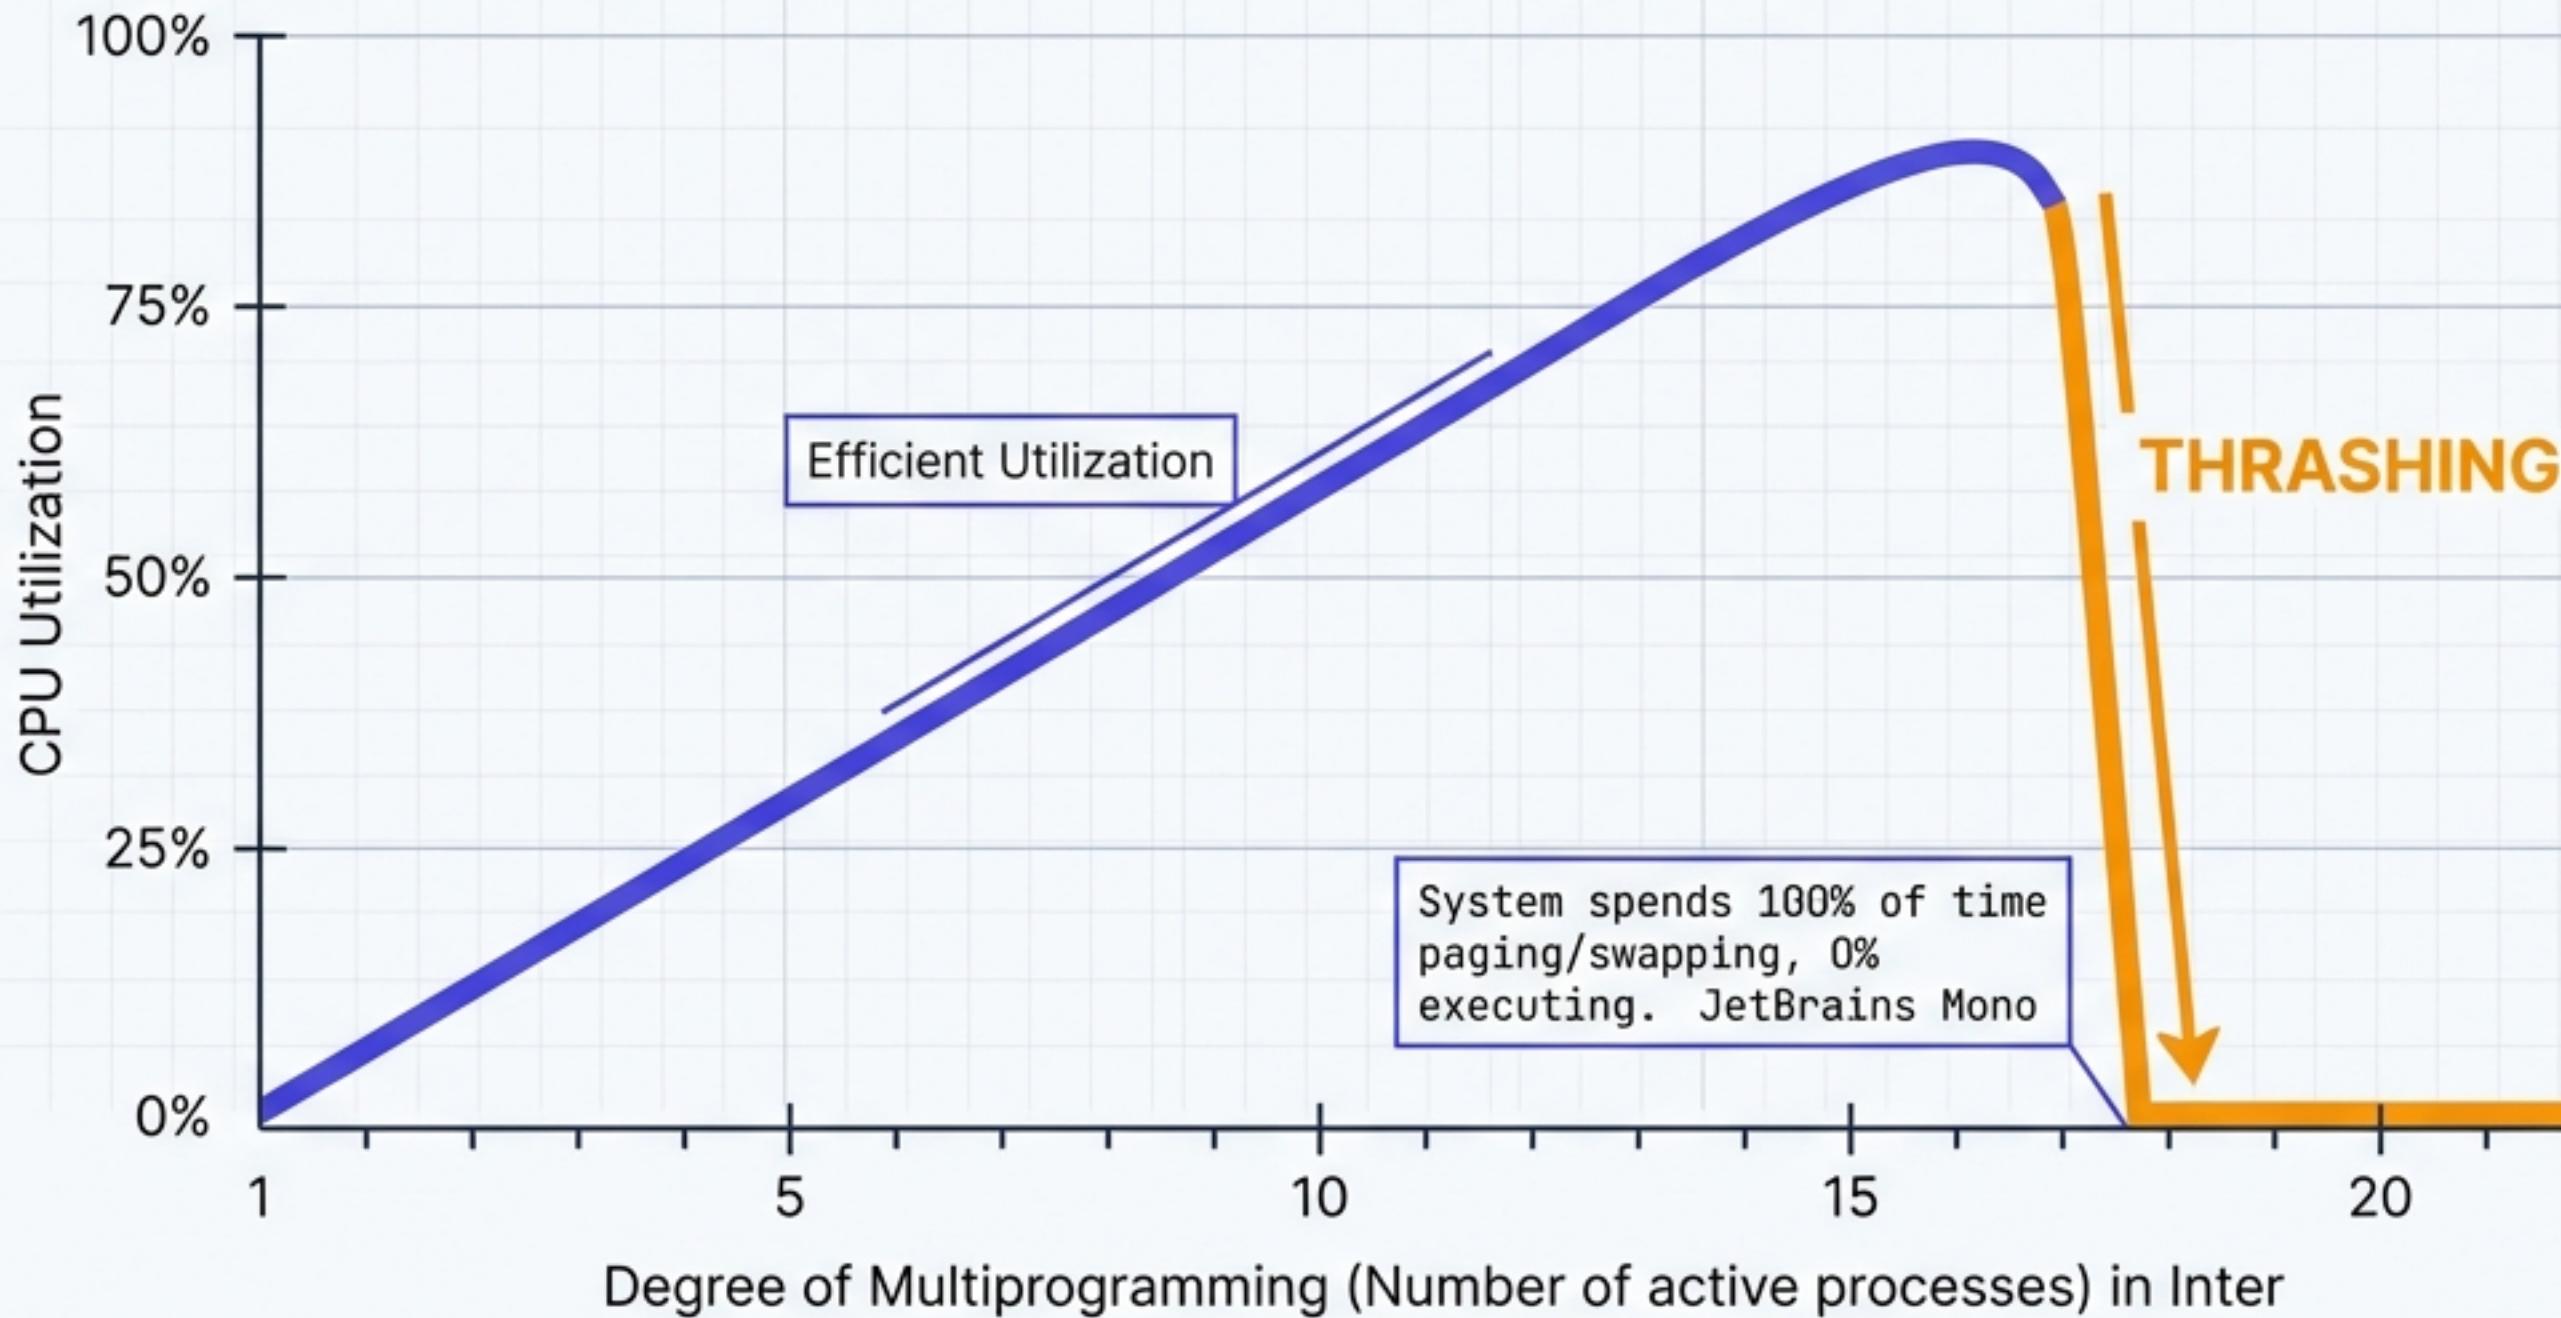


Visible to user logic. Segments vary in size. Causes external fragmentation (holes between segments too small to use).

# Memory Performance & Pitfalls

The crucial balance between process execution and paging overhead.

## The Thrashing Curve



### ❶ Locality of Reference:

Programs tend to access specific clusters of memory repeatedly. Caching relies on this.

### ❷ Belady's Anomaly:

A paradox in FIFO paging where adding *more* memory frames actually causes *more* page faults.

# Storage & I/O Architecture

## RAID Levels (Redundant Array of Independent Disks)

### RAID 0 (Striping)



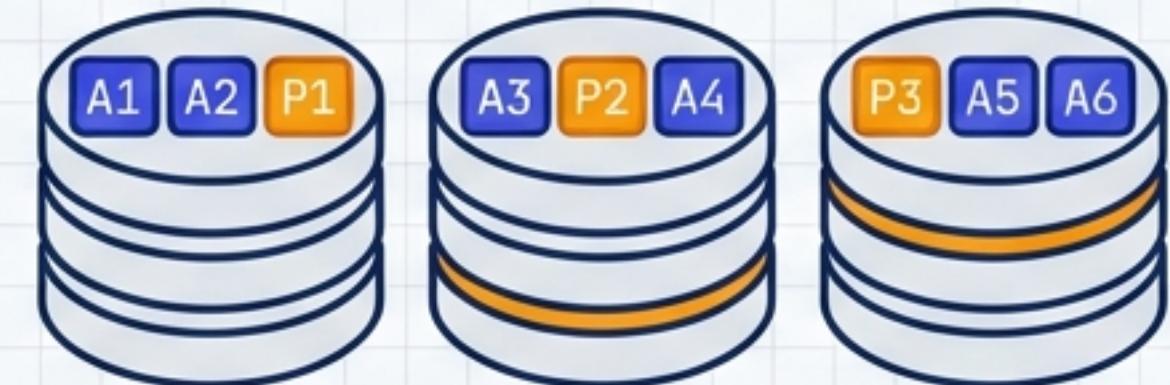
(Fast, No Redundancy)

### RAID 1 (Mirroring)



(Redundant, Expensive)

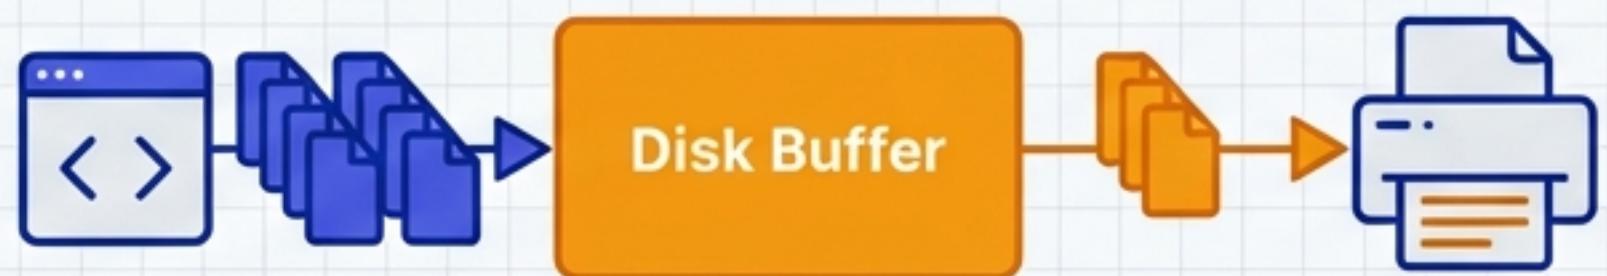
### RAID 5 (Striping + Parity)



(Balanced)

## I/O Techniques

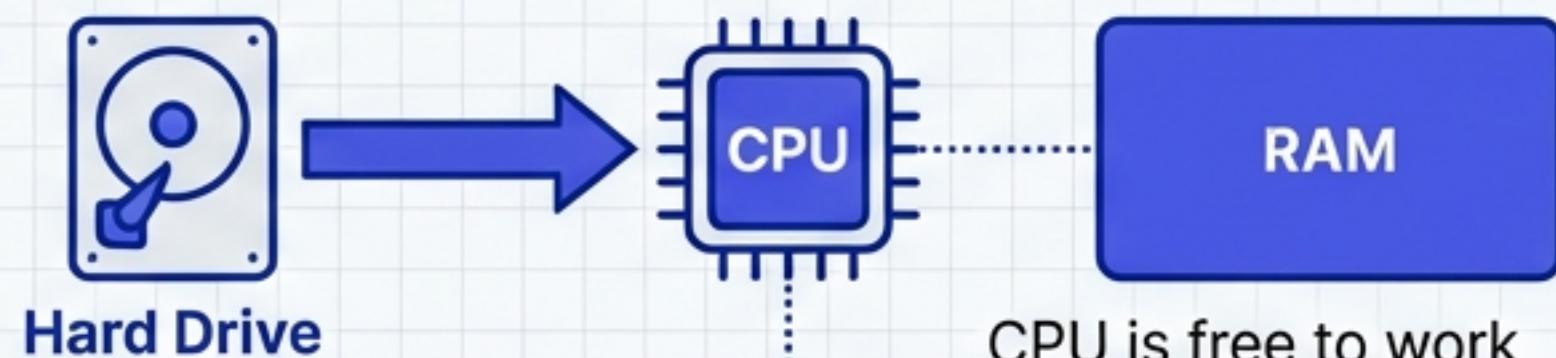
### Spooling



App

Overlap computation with I/O.

### DMA (Direct Memory Access)



CPU is free to work while data moves.

# System Design & Terminology Cheat Sheet

Rapid-fire definitions for interview readiness.

## Interrupt vs. Trap

**Interrupt:** Hardware signal (async) needing attention.

**Trap:** Software interrupt (sync) caused by error or system call.

## Cycle Stealing

I/O controllers accessing RAM via the bus when the CPU isn't using it, subtly 'stealing' memory cycles.

## Compaction

Shuffling memory contents to place all free space together in one large block. Solves external fragmentation.

## Overlay

Loading only the required instructions of a program into memory, overwriting previously used parts. (Old technique).

## Bit-Vector

A string of bits (0010110) used to track free disk blocks.  
1 = Free, 0 = Allocated.

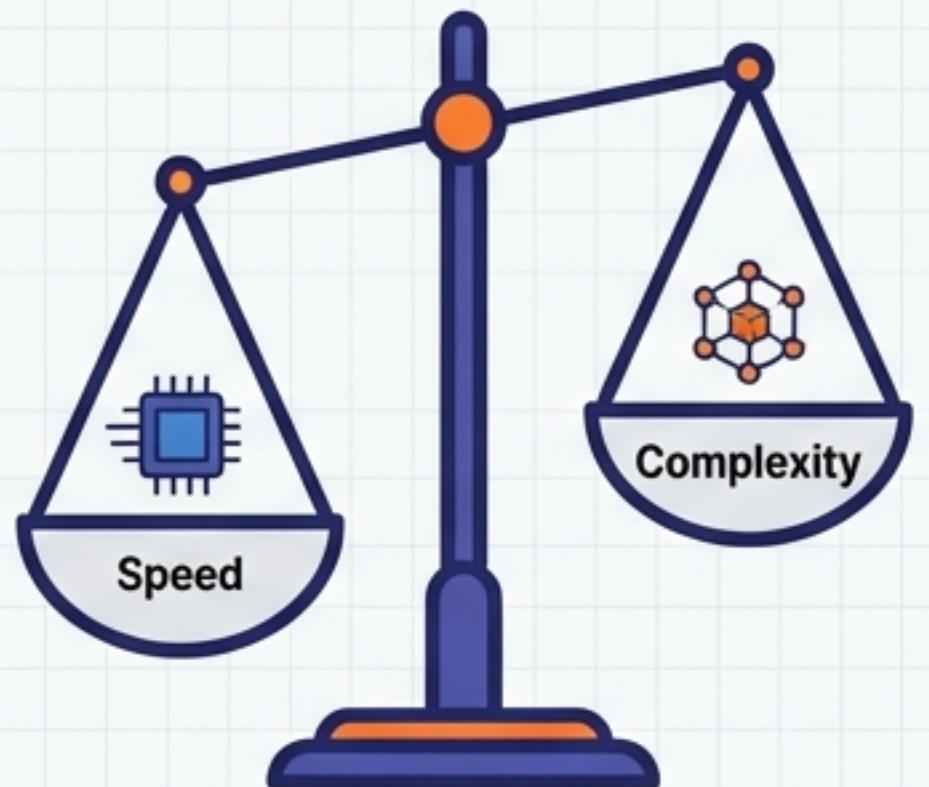
## Starvation vs. Aging

**Starvation:** Low priority process never runs.

**Aging:** Gradually increasing priority over time to ensure execution.

# The Balancing Act: Key Takeaways

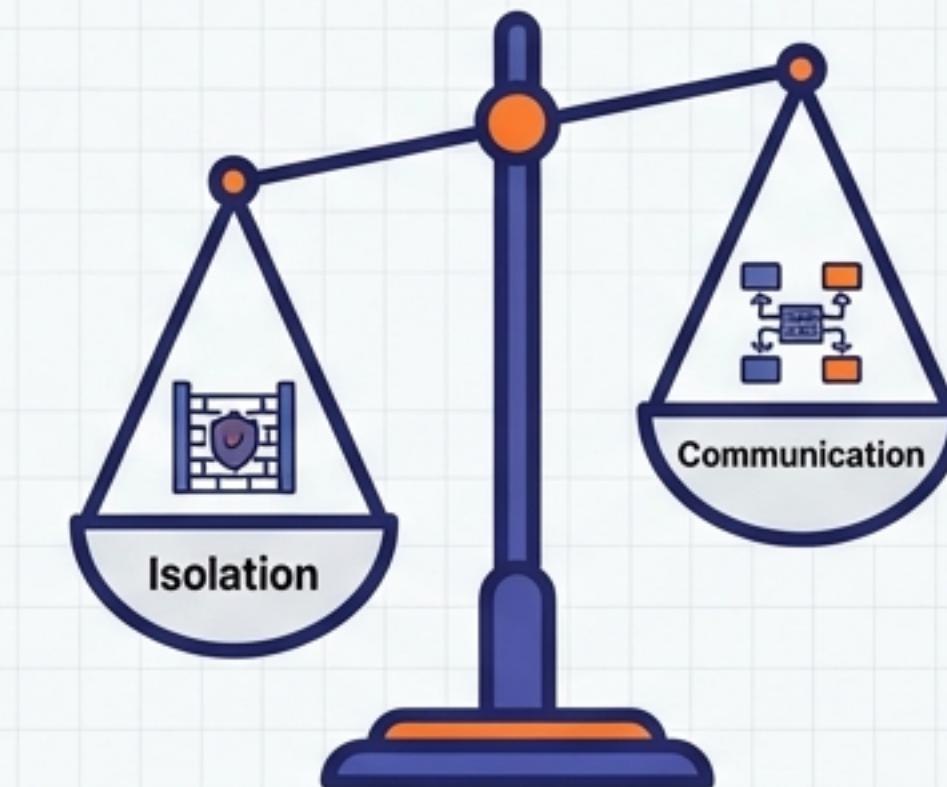
Operating System design is the management of trade-offs.



**Monolithic kernels** offer raw speed but difficult maintenance. **Microkernels** offer stability but suffer **IPC overhead**.



**Preemptive scheduling** is responsive but costly (**context switching**). **Non-preemptive** is efficient but allows **starvation**.



**Processes** provide safety through isolation. **Threads** provide speed through **shared memory** but introduce risk (**Race Conditions**).

*A robust OS manages the chaos of hardware concurrency to provide the illusion of infinite, dedicated resources to the user.*